

▼ S223090226

SHREYAS VIVEK

SIT719 5.2D 2

```
import os
from collections import defaultdict
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, f1_score, precision_score, recall_score, classification_report
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.preprocessing import StandardScaler
import warnings
import seaborn as sns
from time import time
```

```
# %matplotlib inline
warnings.filterwarnings('ignore')

import pandas as pd

import io
```

```
dataset = cdataset.dropna()
```

```
cdataset = pd.read_csv('Processed_Combined_IoT_dataset.csv')
```

```
dataset.head()
```


	FC1_Read_Input_Register	FC2_Read_Discrete_Value	FC3_Read_Holding_Register	FC4_Read_Coil	current_
0	0.495216	0.499092	0.488897	0.499405	
1	0.495216	0.499092	0.488897	0.499405	
2	0.495216	0.499092	0.488897	0.499405	
3	0.495216	0.499092	0.488897	0.499405	
4	0.495216	0.499092	0.488897	0.499405	

```
print(dataset.shape)
```

```
(401119, 18)
```

```
print(list(dataset.columns))
```

```
['FC1_Read_Input_Register', 'FC2_Read_Discrete_Value', 'FC3_Read_Holding_Register', 'FC4_Read_Coil', 'current_temperature', 'door_state']
```



```
target_cols=list(dataset.columns[-1:])  
target_cols
```

```
['label']
```

```
feature_cols= list(dataset.columns[:-1])  
feature_cols
```

```
['FC1_Read_Input_Register',  
 'FC2_Read_Discrete_Value',  
 'FC3_Read_Holding_Register',  
 'FC4_Read_Coil',  
 'current_temperature',  
 'door_state',  
 'fridge_temperature',  
 'humidity',  
 'latitude',  
 'light_status',  
 'longitude',  
 'motion_status',  
 'pressure',  
 'sphone_signal',  
 'temp_condition',  
 'temperature',  
 'thermostat_status']
```

```
#split dataset in features and target variable
X = dataset.drop('label', axis=1) # Features
y = dataset['label'] # Target variable

X.head()
```

	FC1_Read_Input_Register	FC2_Read_Discrete_Value	FC3_Read_Holding_Register	FC4_Read_Coil	current_
0	0.495216	0.499092	0.488897	0.499405	
1	0.495216	0.499092	0.488897	0.499405	
2	0.495216	0.499092	0.488897	0.499405	
3	0.495216	0.499092	0.488897	0.499405	
4	0.495216	0.499092	0.488897	0.499405	

```
y.head()
```

```
0    0
1    0
2    0
3    0
4    0
Name: label, dtype: int64
```

```
# Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1) # 70% training and 30% test
```

```
# Check the shape of all of these
print("X_train shape is : ", X_train.shape)
print("X_test shape is : ", X_test.shape)
```

```
print("y_train shape is : ", y_train.shape)
print("y_test shape is : ", y_test.shape)
```

```
X_train shape is : (280783, 17)
X_test shape is : (120336, 17)
y_train shape is : (280783,)
y_test shape is : (120336,)
```

```
# Define a list of classifier classes
classifiers = [
    DecisionTreeClassifier,
    RandomForestClassifier,
    KNeighborsClassifier,
    LogisticRegression,
    GaussianNB,
]
```

```
def evaluate_classifier(class_type, train_x, train_Y, test_x, test_Y, label_names=None, **kwargs):
    start_time = time() # Record the start time for training

    if class_type == 'DecisionTree':
        from sklearn.tree import DecisionTreeClassifier
        classifier = DecisionTreeClassifier(**kwargs)
    elif class_type == 'RandomForest':
        from sklearn.ensemble import RandomForestClassifier
        classifier = RandomForestClassifier(**kwargs)
    elif class_type == 'KNeighbors':
        from sklearn.neighbors import KNeighborsClassifier
        classifier = KNeighborsClassifier(**kwargs)
    elif class_type == 'LogisticRegression':
        from sklearn.linear_model import LogisticRegression
        classifier = LogisticRegression(**kwargs)
    elif class_type == 'NaiveBayes':
        from sklearn.naive_bayes import GaussianNB
        classifier = GaussianNB(**kwargs)
    else:
        raise ValueError(f"Unsupported classifier type: {class_type}")
```

```

classifier.fit(train_x, train_Y)
training_time = time() - start_time # Calculate training time

start_time = time() # Record the start time for testing
y_pred = classifier.predict(test_x)
testing_time = time() - start_time # Calculate testing time

class_name = class_type
print(f"Classifier: {class_name}")
print(f"Evaluation Report for {class_name}:")
show_evaluation_results(test_Y, y_pred, label_names, training_time, testing_time)
print("=" * 40)

```

```

label_names = list(map(str, np.unique(y_test)))
performance_metrics = []

```

```

def show_evaluation_results(test_Y, y_pred, label_names, training_time, testing_time):
    accuracy = accuracy_score(test_Y, y_pred)
    conf_matx = confusion_matrix(test_Y, y_pred)
    f1score = f1_score(test_Y, y_pred, average="macro")
    precision = precision_score(test_Y, y_pred, average="macro")
    recall = recall_score(test_Y, y_pred, average="macro")

    print(f"F-Score: {f1score}")
    print(f"Precision: {precision}")
    print(f"Re-call: {recall}")
    print(f"Accuracy: {accuracy}")
    print(f"Confusion Matrix:\n{conf_matx}")

    clrp = classification_report(test_Y, y_pred, target_names=label_names)
    print(clrp)

    class_far = calculate_false_alarm_rate(conf_matx, label_names)
    for label_name, false_alarm in class_far.items():
        print(f"False Alarm of {label_name}: {false_alarm:.4f} ({false_alarm * 100:.2f}%)")

    overall_far = sum(class_far.values()) / len(class_far)
    print(f"Overall False Alarm Rate: {overall_far:.4f} ({overall_far * 100:.2f}%)")

```

```
# Add training time, testing time, and error rate to the report
print(f"Training Time: {training_time:.2f} seconds")
print(f"Testing Time: {testing_time:.2f} seconds")
error_rate = 1 - accuracy
print(f"Error Rate: {error_rate:.4f} ({error_rate * 100:.2f}%)")

# Create the heatmap using Matplotlib
plt.imshow(conf_matx, interpolation='nearest', cmap='plasma')
plt.title("Confusion Matrix")
plt.colorbar()

# Label the axes with numbers
for i in range(len(label_names)):
    for j in range(len(label_names)):
        plt.text(j, i, str(conf_matx[i, j]), ha='center', va='center', color='white')

# Label the axes
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.xticks(np.arange(len(label_names)), label_names, rotation=45)
plt.yticks(np.arange(len(label_names)), label_names)

# Append performance metrics to the list
performance_metrics.append({
    'F-Score': f1score,
    'Precision': precision,
    'Recall': recall,
    'Accuracy': accuracy,
    'Overall FAR': overall_far,
    'Training Time': training_time,
    'Testing Time': testing_time,
    'Error Rate': error_rate
})

plt.show()
```

```
def calculate_false_alarm_rate(conf_matx, label_names):
    class_far = {}
```

```
for i, label_name in enumerate(label_names):
    TP = conf_matx[i, i]
    FP = np.sum(conf_matx[:, i]) - TP
    FN = np.sum(conf_matx[i, :]) - TP
    TN = np.sum(conf_matx) - TP - FP - FN
    false_alarm = FP / (FP + TN)
    class_far[label_name] = false_alarm
return class_far
```

```
classifiers = [
    ('DecisionTree', DecisionTreeClassifier(random_state=17)),
    ('RandomForest', RandomForestClassifier(random_state=17)),
    ('KNeighbors', KNeighborsClassifier()),
    ('LogisticRegression', LogisticRegression()),
    ('NaiveBayes', GaussianNB()),
]
```

```
for classifier_name, _ in classifiers:
    evaluate_classifier(class_type=classifier_name, train_x=X_train, train_Y=y_train, test_x=X_test, test_Y=y_test, label_names=label_names)
metrics_names = ['F-Score', 'Precision', 'Recall', 'Accuracy', 'Overall FAR', 'Training Time', 'Testing Time', 'Error Rate']
metrics_values = [list(metric.values()) for metric in performance_metrics]

plt.figure(figsize=(12, 12)) # Increase the figure size to accommodate additional charts
```




=====

Classifier: KNeighbors

Evaluation Report for KNeighbors:

F-Score: 0.8188954292131032

Precision: 0.8471684784464781

Re-call: 0.8072729637735185

Accuracy: 0.8362667863316049

Confusion Matrix:

[[68951 4544]

[15159 31682]]

	precision	recall	f1-score	support
0	0.82	0.94	0.87	73495
1	0.87	0.68	0.76	46841
accuracy			0.84	120336
macro avg	0.85	0.81	0.82	120336
weighted avg	0.84	0.84	0.83	120336

False Alarm of 0: 0.3236 (32.36%)

False Alarm of 1: 0.0618 (6.18%)


Overall False Alarm Rate: 0.1927 (19.27%)

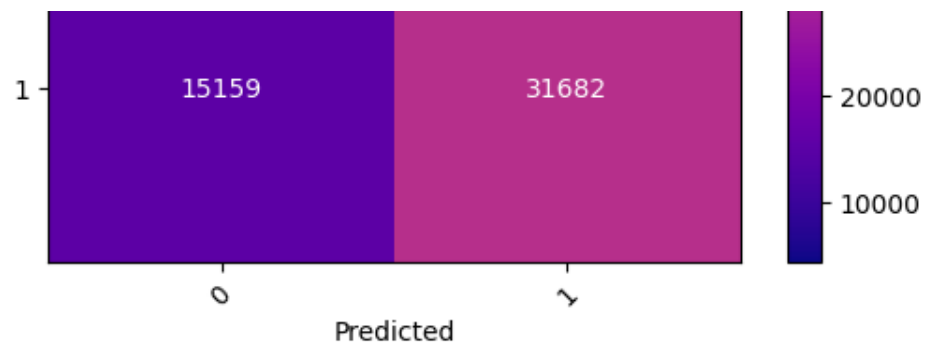
Training Time: 0.05 seconds

Testing Time: 184.40 seconds

Error Rate: 0.1637 (16.37%)

Confusion Matrix





=====

Classifier: LogisticRegression

Evaluation Report for LogisticRegression:

F-Score: 0.582669930370106

Precision: 0.7631321356723366

Re-call: 0.606567843050947

Accuracy: 0.6881814253423747

Confusion Matrix:

[[71660 1835]

[35688 11153]]

	precision	recall	f1-score	support
0	0.67	0.98	0.79	73495
1	0.86	0.24	0.37	46841
accuracy			0.69	120336
macro avg	0.76	0.61	0.58	120336

