

▼ Importing libraries

```
import os
from collections import defaultdict
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, f1_score, precision_score, recall_score, classification_report
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.preprocessing import StandardScaler
import warnings
import seaborn as sns
from time import time
```

```
%matplotlib inline
warnings.filterwarnings('ignore')
```

```
import pandas as pd
```

```
import io
```

```
dataset = pd.read_csv('Processed_Combined_IoT_dataset.csv')
```

```
dataset.head()
```

	FC1_Read_Input_Register	FC2_Read_Discrete_Value	FC3_Read_Holding_Register	FC4_Read_Coil	current_
0	0.495216	0.499092	0.488897	0.499405	
1	0.495216	0.499092	0.488897	0.499405	
2	0.495216	0.499092	0.488897	0.499405	
3	0.495216	0.499092	0.488897	0.499405	
4	0.495216	0.499092	0.488897	0.499405	

```
print(dataset.shape)
```

```
(401119, 18)
```

```
print(list(dataset.columns))
```

```
['FC1_Read_Input_Register', 'FC2_Read_Discrete_Value', 'FC3_Read_Holding_Register', 'FC4_Read_Coil', 'current_temperature', 'door_s
```

```
target_cols=list(dataset.columns[-1:])
target_cols
```

```
['label']
```

```
feature_cols= list(dataset.columns[:-1])
feature_cols
```

```
['FC1_Read_Input_Register',
 'FC2_Read_Discrete_Value',
 'FC3_Read_Holding_Register',
 'FC4_Read_Coil',
 'current_temperature',
 'door_state',
 'fridge_temperature',
 'humidity',
 'latitude',
 'light_status',
 'longitude',
 'motion_status',
 'pressure',
 'sphone_signal',
```

```
'temp_condition',
'temperature',
'thermostat_status']
```

```
#split dataset in features and target variable
X = dataset.drop('label', axis=1) # Features
y = dataset['label'] # Target variable
```

```
X.head()
```

	FC1_Read_Input_Register	FC2_Read_Discrete_Value	FC3_Read_Holding_Register	FC4_Read_Coil	current_
0	0.495216	0.499092	0.488897	0.499405	
1	0.495216	0.499092	0.488897	0.499405	
2	0.495216	0.499092	0.488897	0.499405	
3	0.495216	0.499092	0.488897	0.499405	
4	0.495216	0.499092	0.488897	0.499405	

```
y.head()
```

```
0    0
1    0
2    0
3    0
4    0
Name: label, dtype: int64
```

```
# Split dataset into training set and test set
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1) # 70% training and 30% test
```

```
# Check the shape of all of these
```

```
print("X_train shape is : ", X_train.shape)
print("X_test shape is : ", X_test.shape)
print("y_train shape is : ", y_train.shape)
print("y_test shape is : ", y_test.shape)
```

```
X_train shape is : (280783, 17)
X_test shape is : (120336, 17)
y_train shape is : (280783,)
y_test shape is : (120336,)
```

```
# Define a list of classifier classes
```

```
classifiers = [
    DecisionTreeClassifier,
    RandomForestClassifier,
    KNeighborsClassifier,
    LogisticRegression,
    GaussianNB,
]
```

```
def evaluate_classifier(class_type, train_x, train_Y, test_x, test_Y, label_names=None, **kwargs):
```

```
    start_time = time() # Record the start time for training
```

```
    if class_type == 'DecisionTree':
```

```
        from sklearn.tree import DecisionTreeClassifier
        classifier = DecisionTreeClassifier(**kwargs)
```

```
    elif class_type == 'RandomForest':
```

```
        from sklearn.ensemble import RandomForestClassifier
        classifier = RandomForestClassifier(**kwargs)
```

```
    elif class_type == 'KNeighbors':
```

```
        from sklearn.neighbors import KNeighborsClassifier
        classifier = KNeighborsClassifier(**kwargs)
```

```
    elif class_type == 'LogisticRegression':
```

```
        from sklearn.linear_model import LogisticRegression
        classifier = LogisticRegression(**kwargs)
```

```
    elif class_type == 'NaiveBayes':
```

```
        from sklearn.naive_bayes import GaussianNB
        classifier = GaussianNB(**kwargs)
```

```
    else:
```

```
        raise ValueError(f"Unsupported classifier type: {class_type}")
```

```
    classifier.fit(train_x, train_Y)
```

```
    training_time = time() - start_time # Calculate training time
```

```
    start_time = time() # Record the start time for testing
```

```

y_pred = classifier.predict(test_x)
testing_time = time() - start_time # Calculate testing time

class_name = class_type
print(f"Classifier: {class_name}")
print(f"Evaluation Report for {class_name}:")
show_evaluation_results(test_Y, y_pred, label_names, training_time, testing_time)
print("=" * 40)

label_names = list(map(str, np.unique(y_test)))
performance_metrics = []

def show_evaluation_results(test_Y, y_pred, label_names, training_time, testing_time):
    accuracy = accuracy_score(test_Y, y_pred)
    conf_matx = confusion_matrix(test_Y, y_pred)
    f1score = f1_score(test_Y, y_pred, average="macro")
    precision = precision_score(test_Y, y_pred, average="macro")
    recall = recall_score(test_Y, y_pred, average="macro")

    print(f"F-Score: {f1score}")
    print(f"Precision: {precision}")
    print(f"Re-call: {recall}")
    print(f"Accuracy: {accuracy}")
    print(f"Confusion Matrix:\n{conf_matx}")

    clrp = classification_report(test_Y, y_pred, target_names=label_names)
    print(clrp)

    class_far = calculate_false_alarm_rate(conf_matx, label_names)
    for label_name, false_alarm in class_far.items():
        print(f"False Alarm of {label_name}: {false_alarm:.4f} ({false_alarm * 100:.2f}%)")

    overall_far = sum(class_far.values()) / len(class_far)
    print(f"Overall False Alarm Rate: {overall_far:.4f} ({overall_far * 100:.2f}%)")

    # Add training time, testing time, and error rate to the report
    print(f"Training Time: {training_time:.2f} seconds")
    print(f"Testing Time: {testing_time:.2f} seconds")
    error_rate = 1 - accuracy
    print(f"Error Rate: {error_rate:.4f} ({error_rate * 100:.2f}%)")

    # Create the heatmap using Matplotlib
    plt.imshow(conf_matx, interpolation='nearest', cmap='plasma')
    plt.title("Confusion Matrix")
    plt.colorbar()

    # Label the axes with numbers
    for i in range(len(label_names)):
        for j in range(len(label_names)):
            plt.text(j, i, str(conf_matx[i, j]), ha='center', va='center', color='white')

    # Label the axes
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.xticks(np.arange(len(label_names)), label_names, rotation=45)
    plt.yticks(np.arange(len(label_names)), label_names)

    # Append performance metrics to the list
    performance_metrics.append({
        'F-Score': f1score,
        'Precision': precision,
        'Recall': recall,
        'Accuracy': accuracy,
        'Overall FAR': overall_far,
        'Training Time': training_time,
        'Testing Time': testing_time,
        'Error Rate': error_rate
    })

plt.show()

def calculate_false_alarm_rate(conf_matx, label_names):
    class_far = {}
    for i, label_name in enumerate(label_names):
        TP = conf_matx[i, i]
        FP = np.sum(conf_matx[:, i]) - TP
        FN = np.sum(conf_matx[i, :]) - TP
        TN = np.sum(conf_matx) - TP - FP - FN
        false_alarm = FP / (FP + TN)
        class_far[label_name] = false_alarm

```

```
    return class_far


classifiers = [
    ('DecisionTree', DecisionTreeClassifier(random_state=17)),
    ('RandomForest', RandomForestClassifier(random_state=17)),
    ('KNeighbors', KNeighborsClassifier()),
    ('LogisticRegression', LogisticRegression()),
    ('NaiveBayes', GaussianNB()),
]

for classifier_name, _ in classifiers:
    evaluate_classifier(class_type=classifier_name, train_x=X_train, train_Y=y_train, test_x=X_test, test_Y=y_test, label_names=label_names)
```

Classifier: DecisionTree
Evaluation Report for DecisionTree:
F-Score: 0.8525781297206456
Precision: 0.8724133244340502
Re-call: 0.8423165766178815
Accuracy: 0.8648201701901342
Confusion Matrix:
[[69373 4122]
 [12145 34696]]

	precision	recall	f1-score	support
0	0.85	0.94	0.90	73495
1	0.89	0.74	0.81	46841
accuracy			0.86	120336
macro avg	0.87	0.84	0.85	120336
weighted avg	0.87	0.86	0.86	120336


False Alarm of 0: 0.2593 (25.93%)
False Alarm of 1: 0.0561 (5.61%)
Overall False Alarm Rate: 0.1577 (15.77%)
Training Time: 3.53 seconds
Testing Time: 0.06 seconds
Error Rate: 0.1352 (13.52%)




=====
Classifier: RandomForest
Evaluation Report for RandomForest:
F-Score: 0.8578706200106176
Precision: 0.8797911526372347
Re-call: 0.8468453020774936
Accuracy: 0.8699973407791517
Confusion Matrix:
[[69921 3574]
 [12070 34771]]

	precision	recall	f1-score	support
0	0.85	0.95	0.90	73495
1	0.91	0.74	0.82	46841
accuracy			0.87	120336
macro avg	0.88	0.85	0.86	120336
weighted avg	0.87	0.87	0.87	120336


False Alarm of 0: 0.2577 (25.77%)
False Alarm of 1: 0.0486 (4.86%)
Overall False Alarm Rate: 0.1532 (15.32%)
Training Time: 75.57 seconds
Testing Time: 3.82 seconds
Error Rate: 0.1300 (13.00%)





```
=====  
Classifier: KNeighbors  
Evaluation Report for KNeighbors:  
F-Score: 0.8188954292131032  
Precision: 0.8471684784464781  
Re-call: 0.8072729637735185  
Accuracy: 0.8362667863316049  
Confusion Matrix:  
[[68951 4544]  
 [15159 31682]]  
  
      precision    recall  f1-score   support  
  
 0      0.82      0.94      0.87       73495  
 1      0.87      0.68      0.76       46841  
  
 accuracy          0.84       120336  
 macro avg      0.85      0.81      0.82       120336  
weighted avg      0.84      0.84      0.83       120336
```

False Alarm of 0: 0.3236 (32.36%)
False Alarm of 1: 0.0618 (6.18%)
Overall False Alarm Rate: 0.1927 (19.27%)
Training Time: 0.05 seconds
Testing Time: 186.46 seconds
Error Rate: 0.1637 (16.37%)



```
=====  
Classifier: LogisticRegression  
Evaluation Report for LogisticRegression:  
F-Score: 0.582669930370106  
Precision: 0.7631321356723366  
Re-call: 0.606567843050947  
Accuracy: 0.6881814253423747  
Confusion Matrix:  
[[71660 1835]  
 [35688 11153]]  
  
      precision    recall  f1-score   support  
  
 0      0.67      0.98      0.79       73495  
 1      0.86      0.24      0.37       46841  
  
 accuracy          0.69       120336  
 macro avg      0.76      0.61      0.58       120336  
weighted avg      0.74      0.69      0.63       120336
```

False Alarm of 0: 0.7619 (76.19%)
False Alarm of 1: 0.0250 (2.50%)
Overall False Alarm Rate: 0.3934 (39.34%)
Training Time: 3.97 seconds
Testing Time: 0.02 seconds
Error Rate: 0.3118 (31.18%)





```
metrics_names = ['F-Score', 'Precision', 'Recall', 'Accuracy', 'Overall FAR', 'Training Time', 'Testing Time', 'Error Rate']
metrics_values = [list(metric.values()) for metric in performance_metrics]
```

```
plt.figure(figsize=(12, 12)) # Increase the figure size to accommodate additional charts
```

```
for i, metric_name in enumerate(metrics_names):
    plt.subplot(4, 2, i + 1)
```

```
    # Special handling for Time and Error Rate metrics
```

```
    if metric_name in ['Training Time', 'Testing Time', 'Error Rate']:
```

```
        plt.bar([classifier[0] for classifier in classifiers], [values[i] for values in metrics_values])
```

```
    else:
```

```
        plt.plot([classifier[0] for classifier in classifiers], [values[i] for values in metrics_values], marker='o')
```

```
    plt.title(metric_name)
```

```
    plt.xlabel('Classifier')
```

```
    plt.xticks(rotation=45)
```

```
    plt.ylabel(metric_name)
```

```
    plt.grid(True)
```

```
plt.tight_layout()
```

```
plt.show()
```



